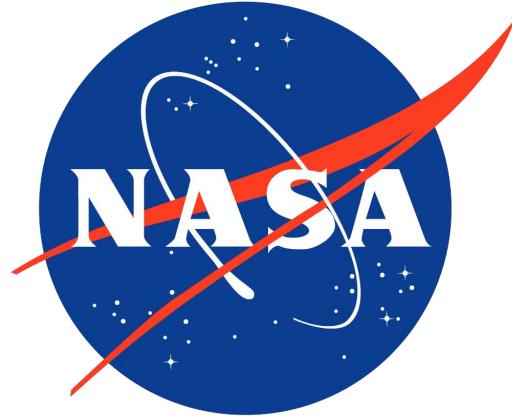




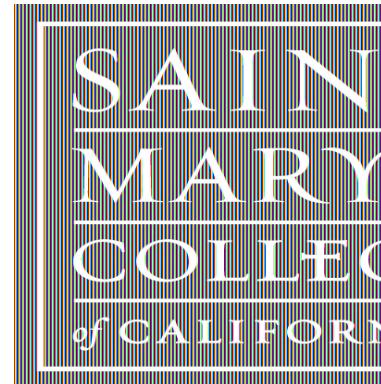
# AUTONOMOUS FLIGHT PLANNING

SAINT MARYS MSBA: S. ZHU K. BRAND C. SUAREZ M. SANTOS

# FOREWORD



We want to thank Nelson, Brian, and Navid for their assistance this project, and all of our former professors throughout the program. Additionally, we want to recognize the support and collaboration we received from other members of the Saint Mary's MSBA program. Thank you!



# AGENDA



"Drones overall will be more impactful than I think people recognize in positive ways to help society."

BILL GATES

- |   |                          |   |             |
|---|--------------------------|---|-------------|
| 1 | Team Introduction        | 6 | Demo        |
| 2 | The Journey of Cohort 12 | 7 | Tech Stack  |
| 3 | Business Problem         | 8 | Future Work |
| 4 | Project Background       | 9 | Questions   |
| 5 | Our Approach             |   |             |

# THE TEAM



## **Kevin Brand**

Saint Mary's College - System Specialist

Hobbies / Interests:

Sports

Cooking

Video Games



## **Colin Suarez**

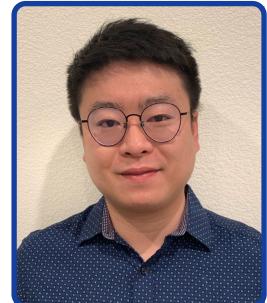
Chiropractic- Office Manager / coaching high school basketball

Hobbies / Interests:

Sports

Graphic design

Travel



## **Steven Zhu**

Walmart.com- Software developer

Hobbies / Interests:

Cooking

Rubiks cubing

Soccer



## **Mark Santos**

Mobility Works- Territory Manager

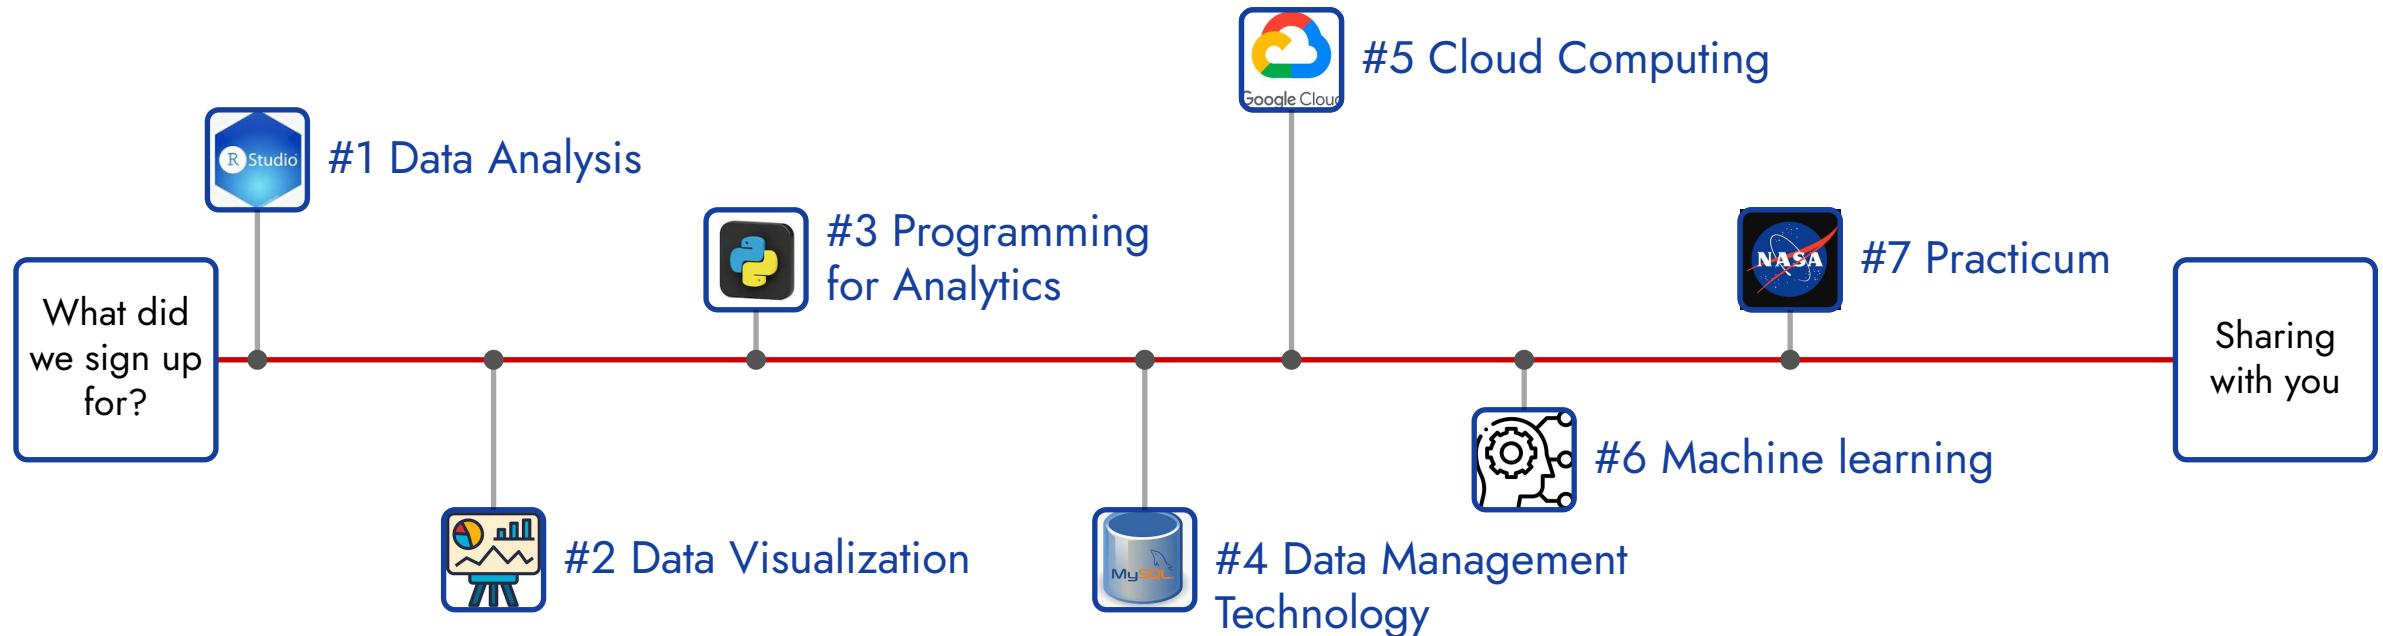
Hobbies / Interests:

Family

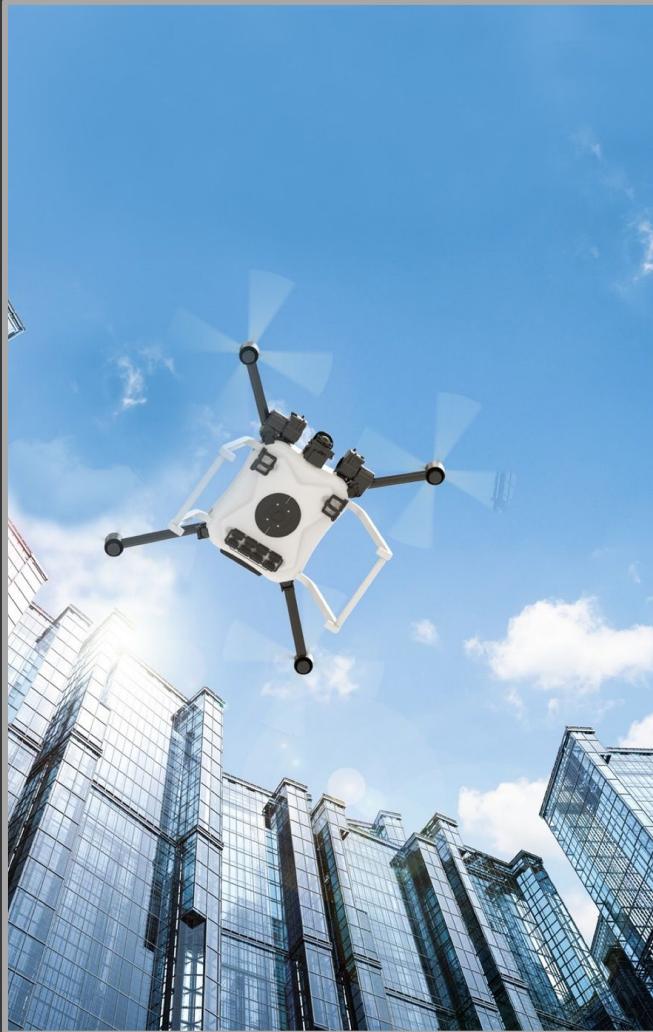
Motorcycles

Mountain Biking

# OUR JOURNEY TO TODAY



# Business Problem



NASA is seeking to build an obstacle detection and avoidance database for unmanned aerial vehicles (UAVs)

This database would allow drones to fly autonomously to their intended destination, saving time and resources. The scope of our project involves building the obstacle database itself in a fast, accurate and efficient way; what the drone does with this information is beyond what we set out to accomplish, but can certainly be explored upon in future work. This presentation will take you through the journey our team took to deliver the requested product.

# PROJECT BACKGROUND

## Why it matters

### Why has it not been solved?

Current technology often has limits on the memory and computation power of most drones. The market needs a solution that can be used efficiently and reliably on ALL drones, not just the high end or military spec drones.

### Why is this a relevant project?

With the growth done use expected to only increase, so does the need for them to be utilized autonomously and safely. For competitive advantages this must be accomplished while ensuring they are using the most optimized flight path to the destination.

# PROJECT BACKGROUND

## Market Research



### Companies using autonomous drones

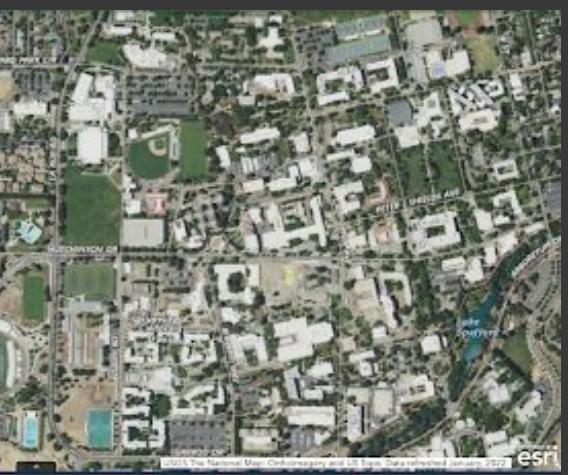
- Amazon
- Google (Wings)
- UPS
- Zipline

### Value Propositions

- Less pollution and traffic from courier vehicles
- Shorter and more accurate ETA for delivery/scheduled delivery
- Minimizing theft of packages
- Increased accuracy of drop offs
- Delivery of urgent supplies

# PROJECT BACKGROUND

Our initial thoughts and questions



## Altitude-

Why can't we just program the drones to fly at an altitude that limits risk of most obstacles?

## Cameras-

Why are we not using camera technology?

## Google-

Why not simply load Google Map view into the drone?

## Satellite Images-

Looking at the USGS satellite images, the closest tile is still very far.  
(8500 feet X 3600 feet Tile) Scale 1:9,028

## LiDAR-

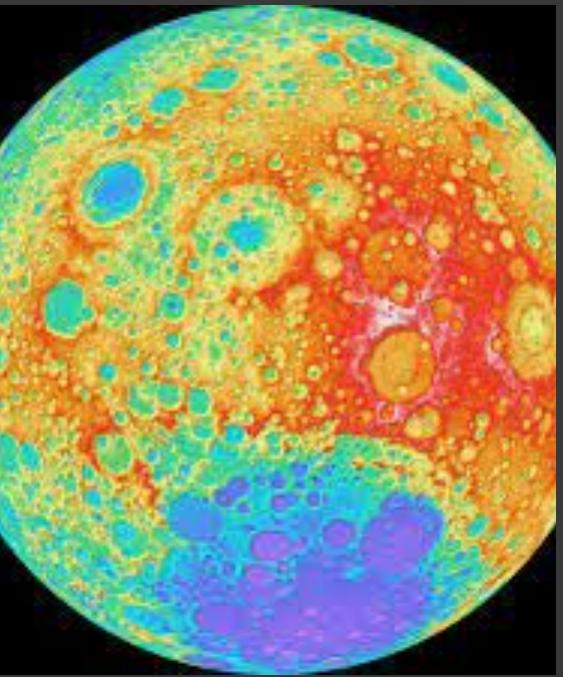
Why are we considering using LiDAR? If Elon says it's wrong for autonomous vehicle usage, it must be wrong for drones, right?

## Data-

Does NASA have access to labeled data?

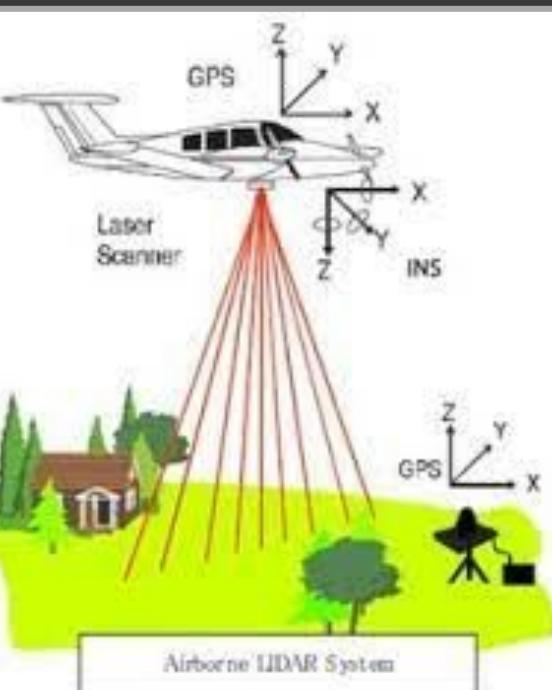
# PROJECT BACKGROUND

## What is LiDAR?



### Development:

LiDAR is an acronym that stands for Light Detection and Ranging. Lidar is a remote sensing method to examine surfaces

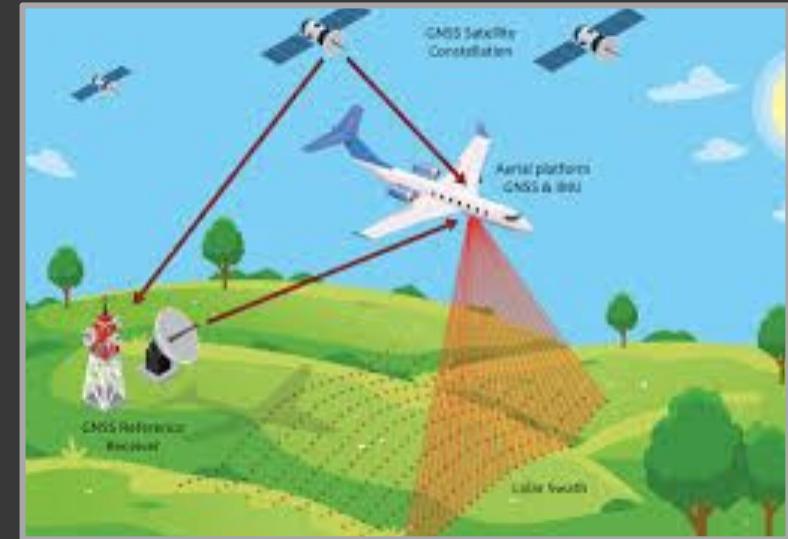
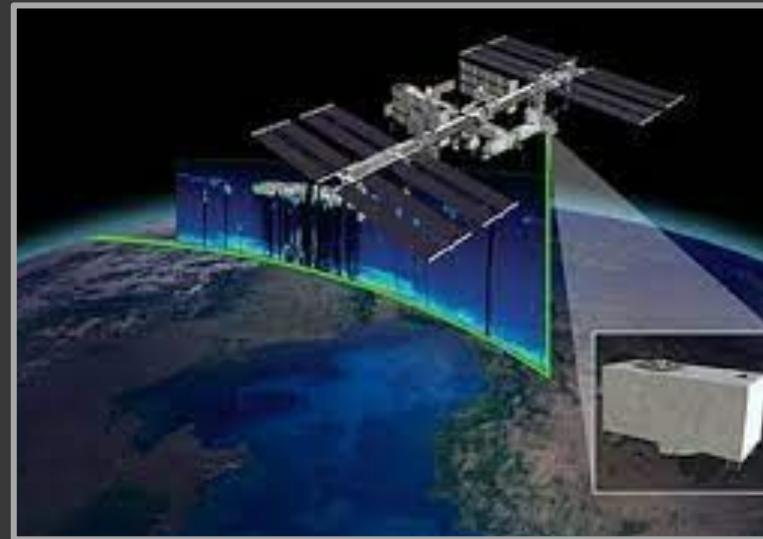
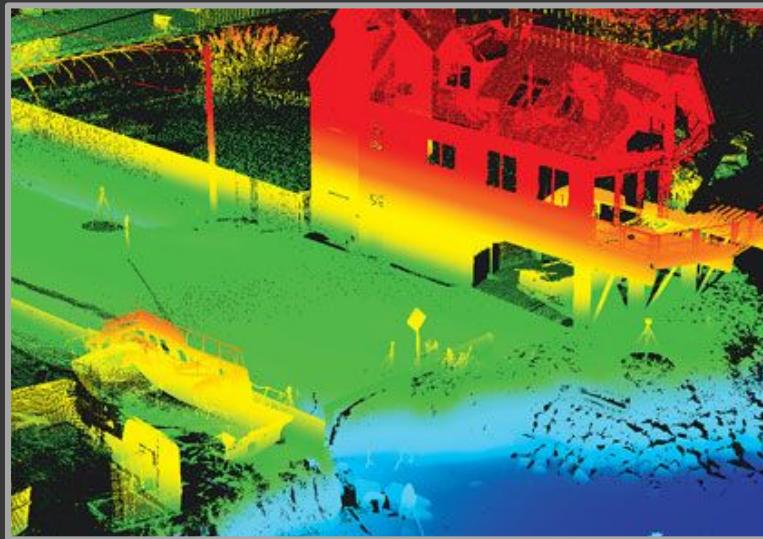


### How does it work:

A LiDAR system calculates how long it takes for beams of light (lasers) to reach an object or a surface and reflect back to the scanner. Depending on the sensors used, LiDAR scanning units can fire hundreds of thousands of pulses per second. By completing this in a quick progression the instrument builds up a complex map of the surface it is measuring

# PROJECT BACKGROUND

## Types of LiDAR



### Terrestrial-

- Scanners are installed on a tripod or vehicle
- Applications in surveying and 3-D models
- Can be Mobile or static

### Satellite-

- LITE Launched in the mid 1990's
- Used to produce atmospheric measurements
- In 2006 CALIPSO was launched

### Aerial-

- Scanners are installed on aircraft
- Most accurate way to create elevation models
- Scans create 3-D point cloud models of landscape

# PROJECT BACKGROUND

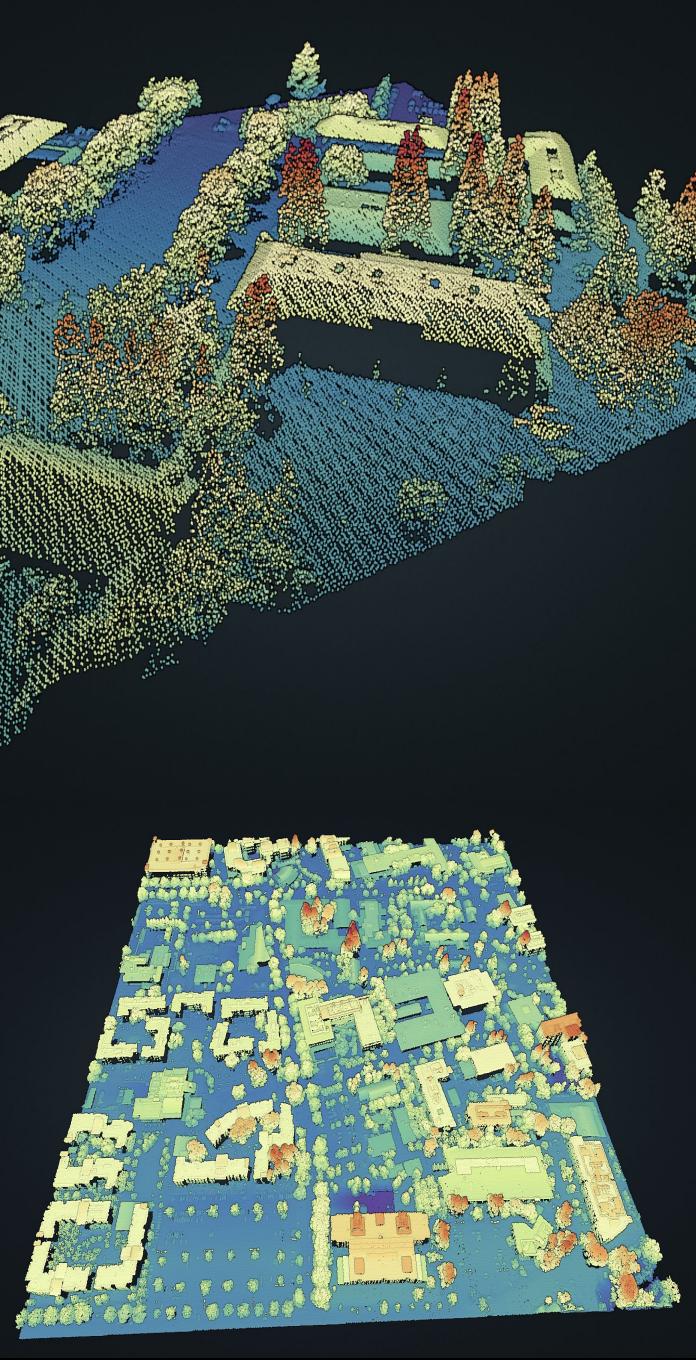
## Point Clouds

### What is a Point Cloud:

Point cloud data is the term used to refer to the data points collected for a given geographical area, terrain, building or space from an aerial scan. A LiDAR point cloud dataset is created when an area is scanned using light detection and ranging

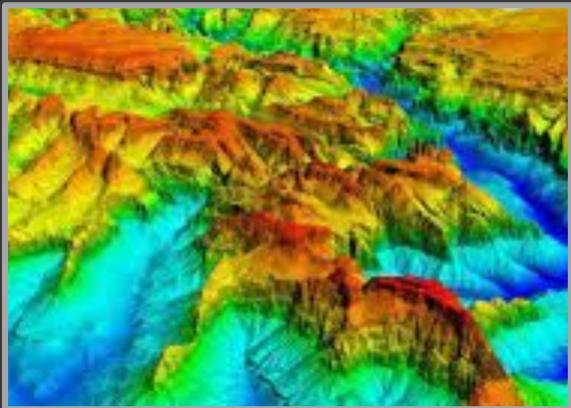
### What is Point Cloud Processing:

Point Cloud processing is a means of turning point cloud data into 3D models



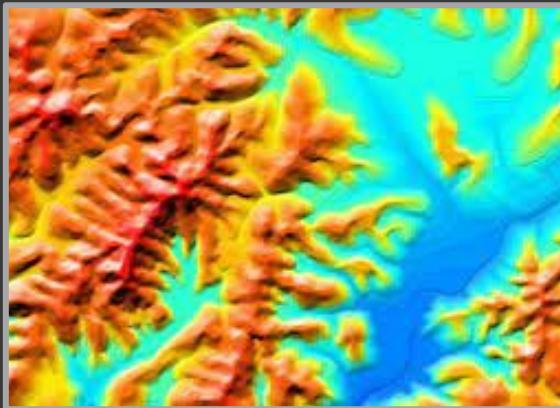
# PROJECT BACKGROUND

## Common LiDAR Data Models



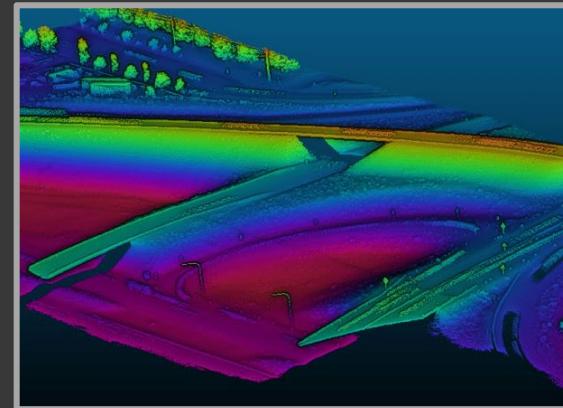
**Digital Elevation Model (DEM)**

Represents the bare-Earth surface, removing all natural and built features



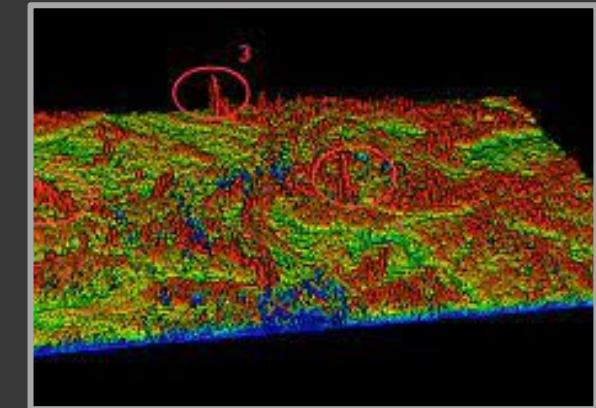
**Digital Terrain Model (DTM)**

Typically augments a DEM, by including vector features of the natural terrain, such as rivers and ridges



**Digital Surface Model (DSM)**

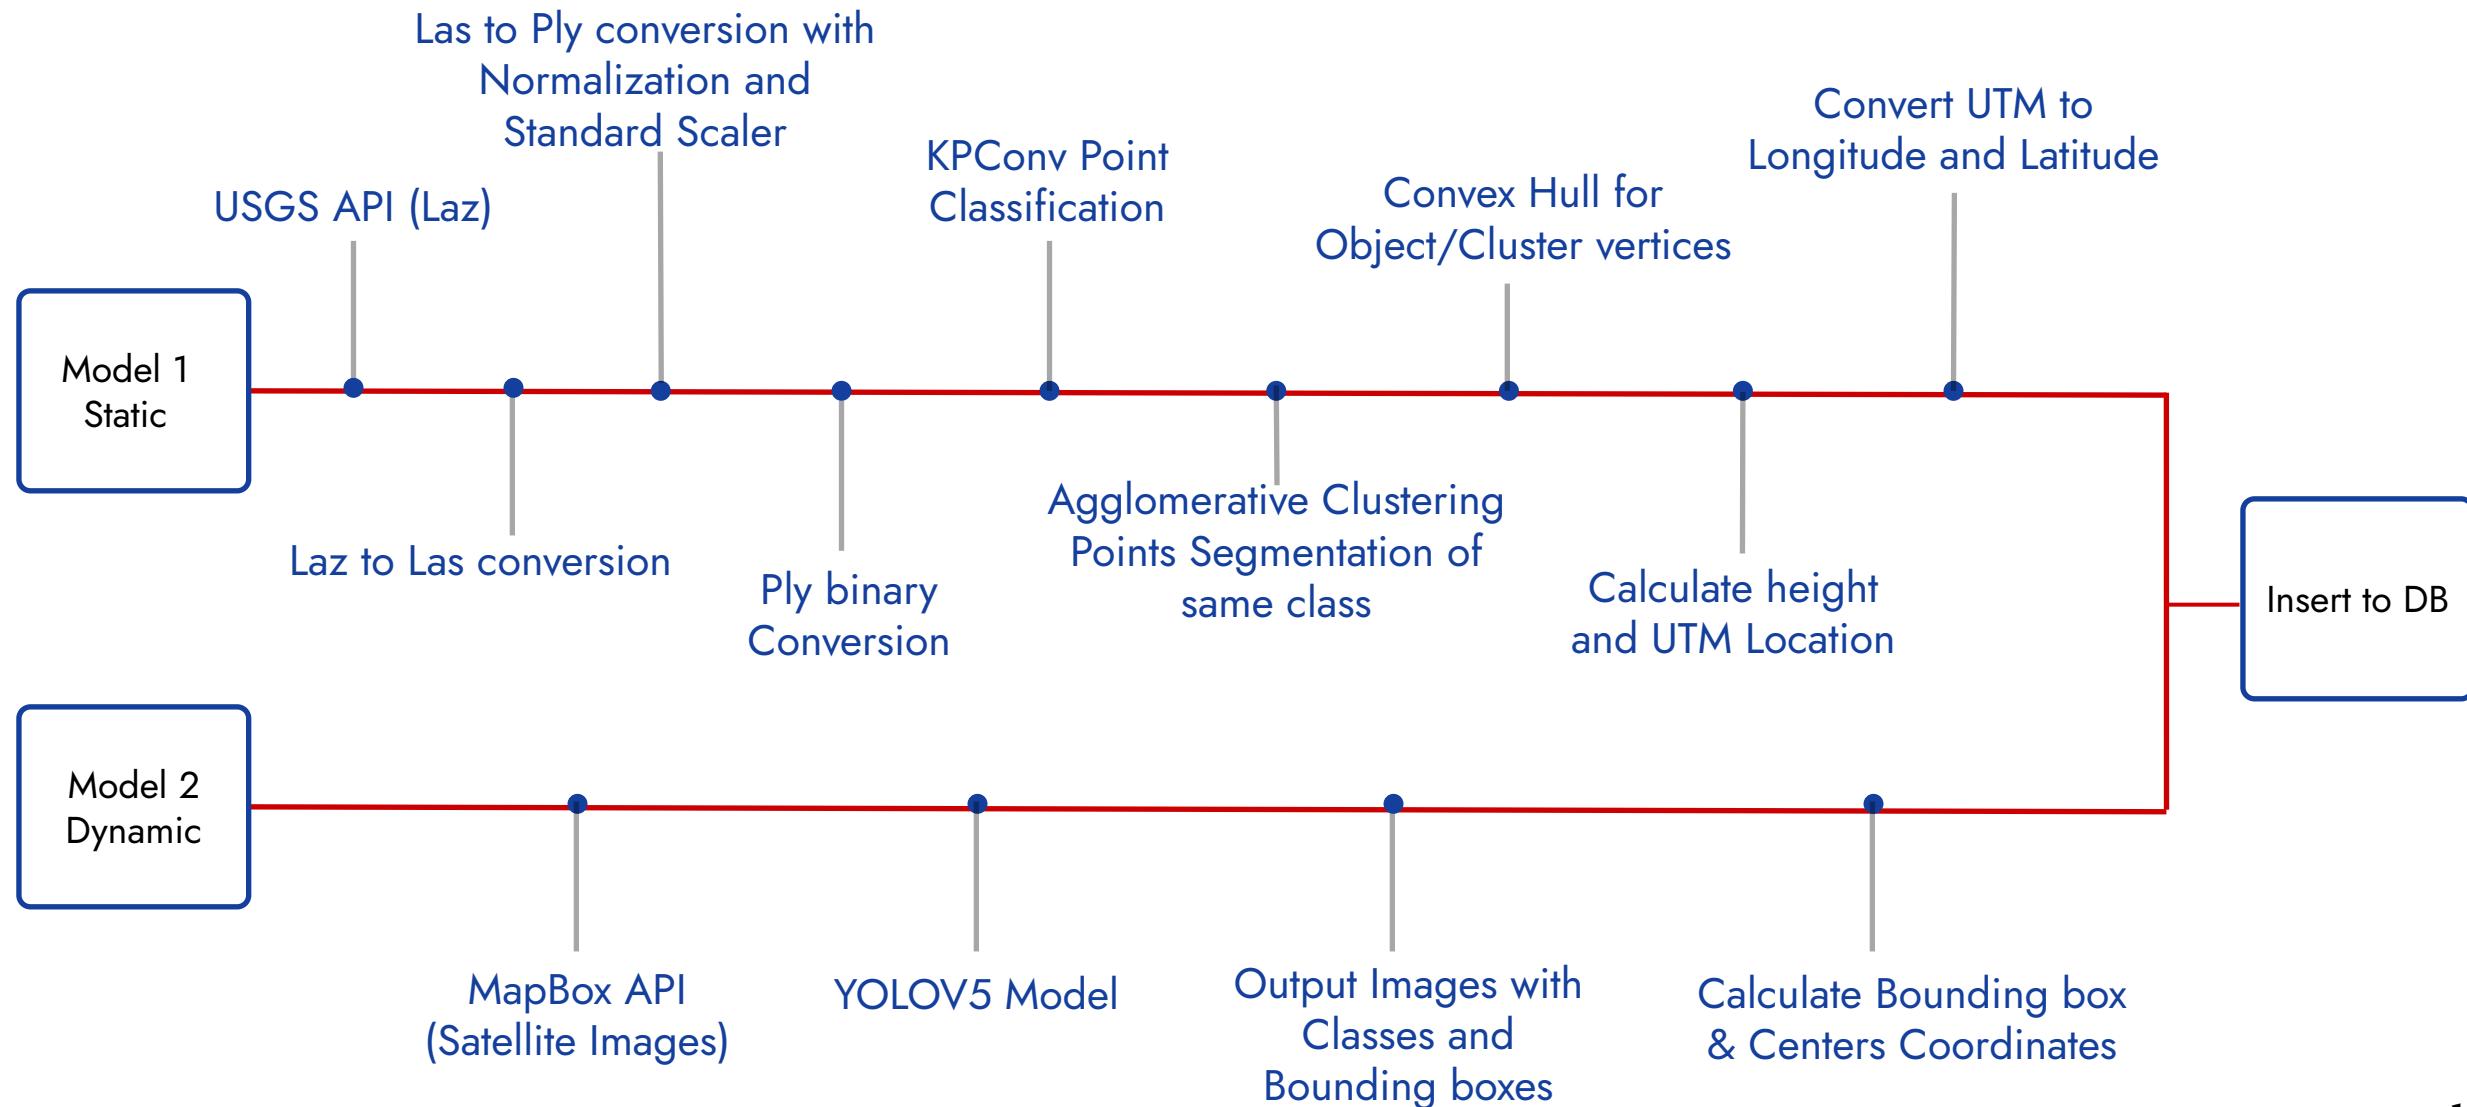
Captures both the natural and built/artificial features of the environment



**Canopy Height Model (CHM)**

Calculated by subtracting the digital terrain model (DTM) from the digital surface model (DSM)

# Overview of Our Approach



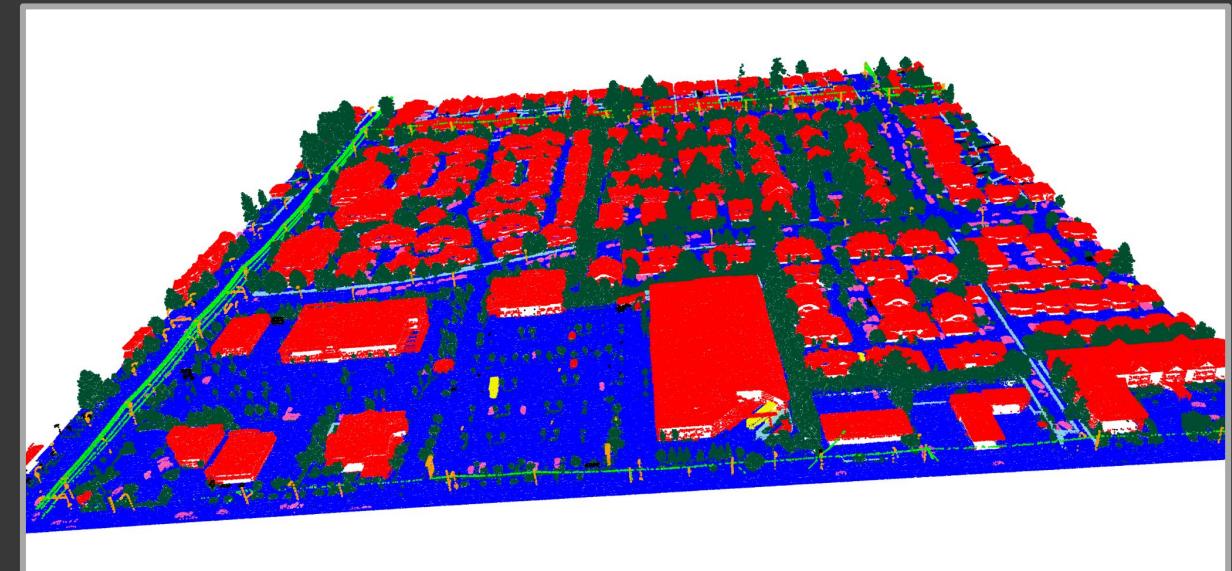
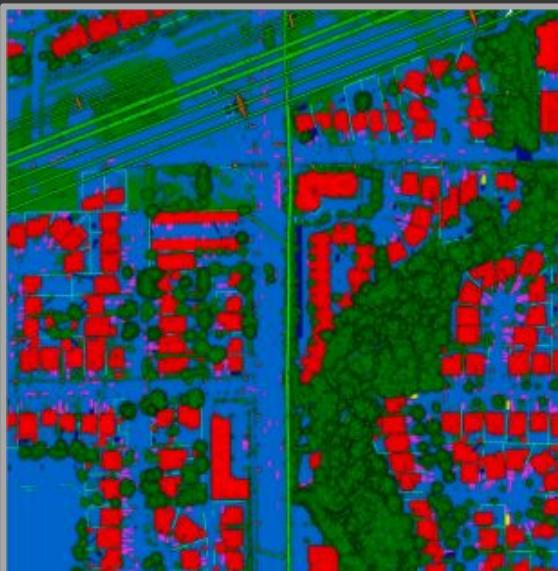
# OUR APPROACH

## The Data

The **DALES** or Dayton Annotated Laser Earth Scan is a large-scale Aerial LiDAR dataset for Point Cloud Segmentation Developed by University of Dayton

Dataset benefits:

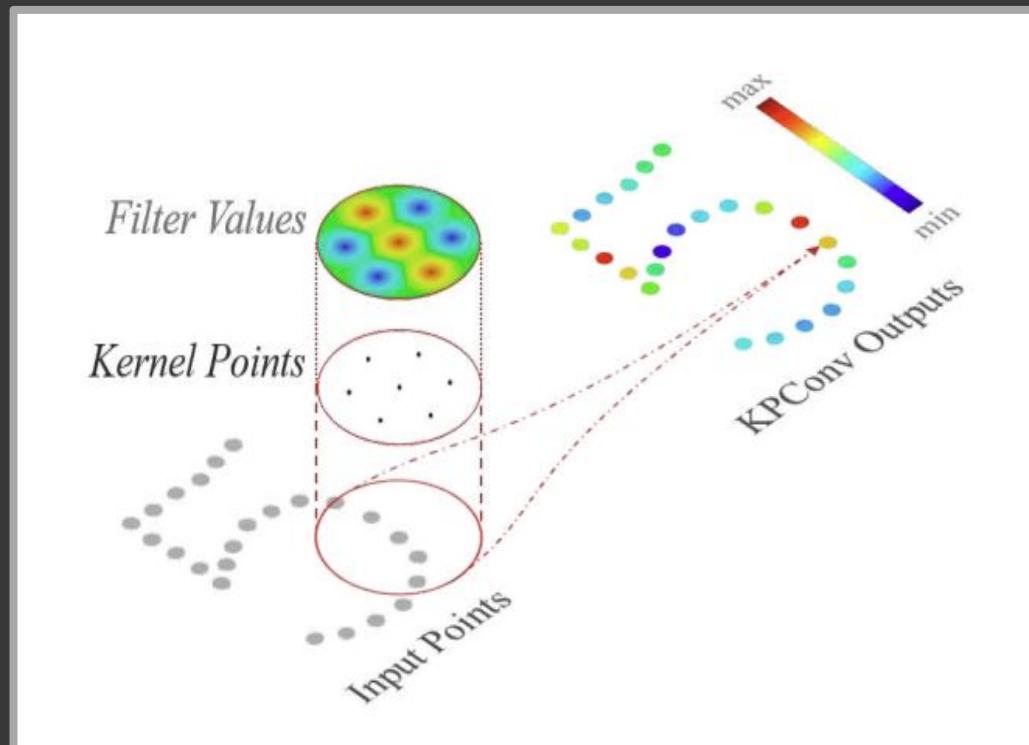
- Contains Over a half billion hand-labeled points
- Pre split training and testing files, cleaned
- Includes two file types: LAS and PLY
- Eight categories: ground, vegetation, cars, trucks, powerlines, poles, fences, and buildings



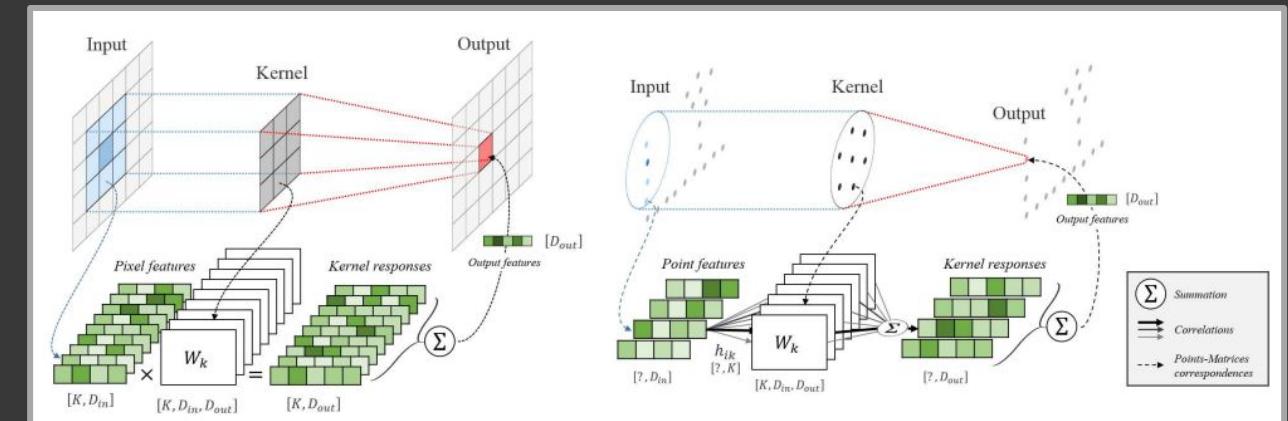
# OUR APPROACH

## First Model - KPConv (Kernel Point Convolution)

- Newer, more advanced point cloud model
- Efficient and robust to varying densities with it's regular subsampling strategy
- Has built in functions for visualization and metrics to help better understand the model



Method	IoU									
	OA	mean	ground	buildings	cars	trucks	poles	power lines	fences	veg
KPConv [8]	<b>0.978</b>	<b>0.811</b>	0.971	<b>0.966</b>	<b>0.853</b>	<b>0.419</b>	<b>0.750</b>	<b>0.955</b>	<b>0.635</b>	<b>0.941</b>
PointNet++ [18]	0.957	0.683	0.941	0.891	0.754	0.303	0.400	0.799	0.462	0.912
ConvPoint [2]	0.972	0.674	0.969	0.963	0.755	0.217	0.403	0.867	0.296	0.919
SuperPoint [10]	0.955	0.606	0.947	0.934	0.629	0.187	0.285	0.652	0.336	0.879
PointCNN [11]	0.972	0.584	<b>0.975</b>	0.957	0.406	0.048	0.576	0.267	0.526	0.917
ShellNet [28]	0.964	0.574	0.960	0.954	0.322	0.396	0.200	0.274	0.600	0.884



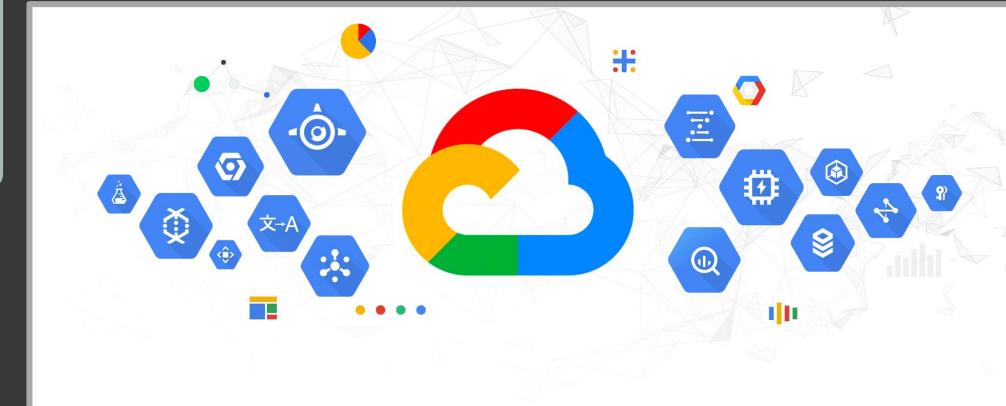
# OUR APPROACH

## Training KPConv on Dales

- o We modified an open source TensorFlow implementation of KPConv to make use of the DALES dataset
- o Using the power of the Google Cloud Platform, we set up a virtual machine to train our model with a stronger GPU and more RAM than our local machines
- o Code and data were loaded onto the VM
- o Training took 6 and a half days to complete a total of 850 epochs
- o The overall model accuracy reached 97.35%

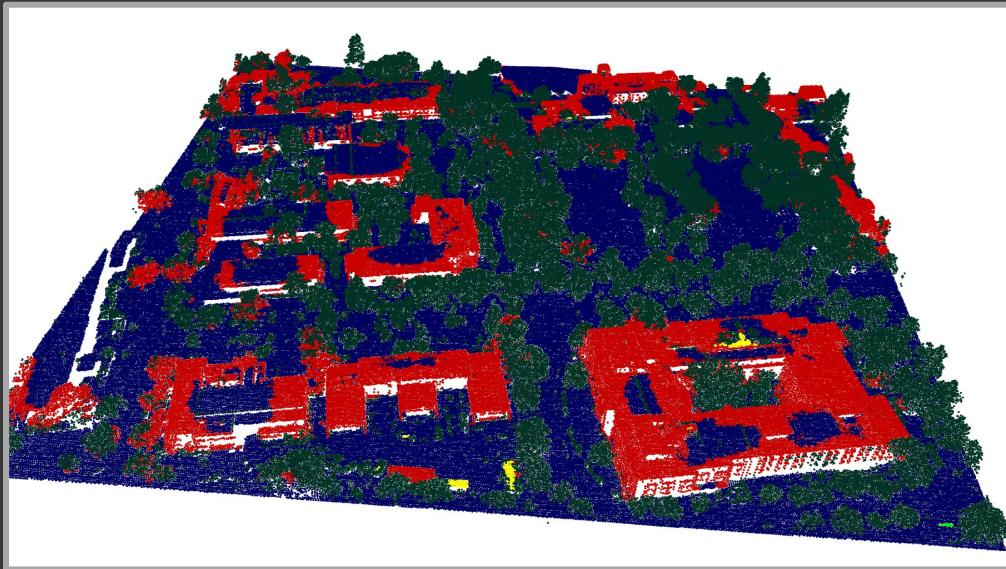
### Machine configuration

Machine type	n1-standard-16
CPU platform	Intel Broadwell
Architecture	x86/64
vCPUs to core ratio <small>?</small>	–
Custom visible cores <small>?</small>	–
Display device	Disabled Enable to use screen capturing and recording tools
GPUs	1 x NVIDIA T4 Virtual Workstation

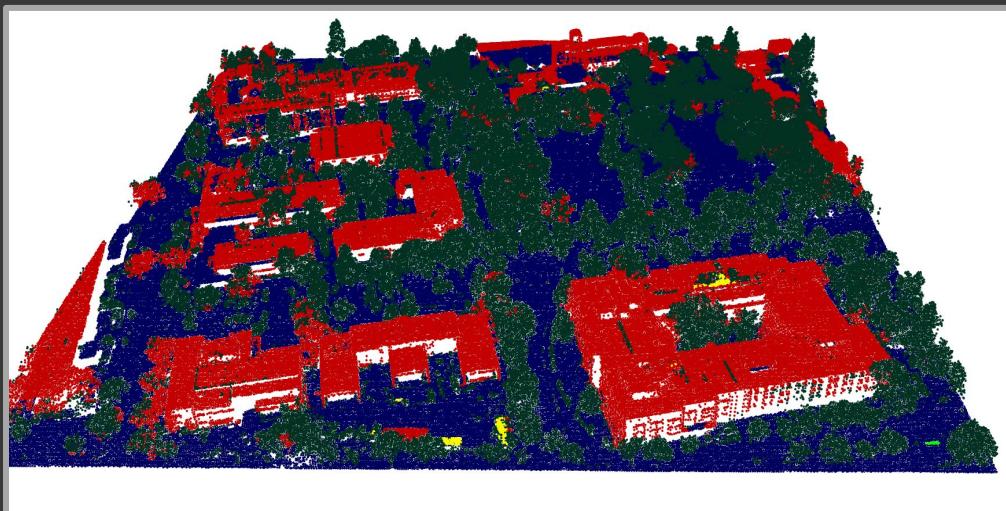


# OUR APPROACH

Testing Model on USGS Data



Our testing on USGS data resulted in buildings classified as ground in some cases, especially if the roof was flat

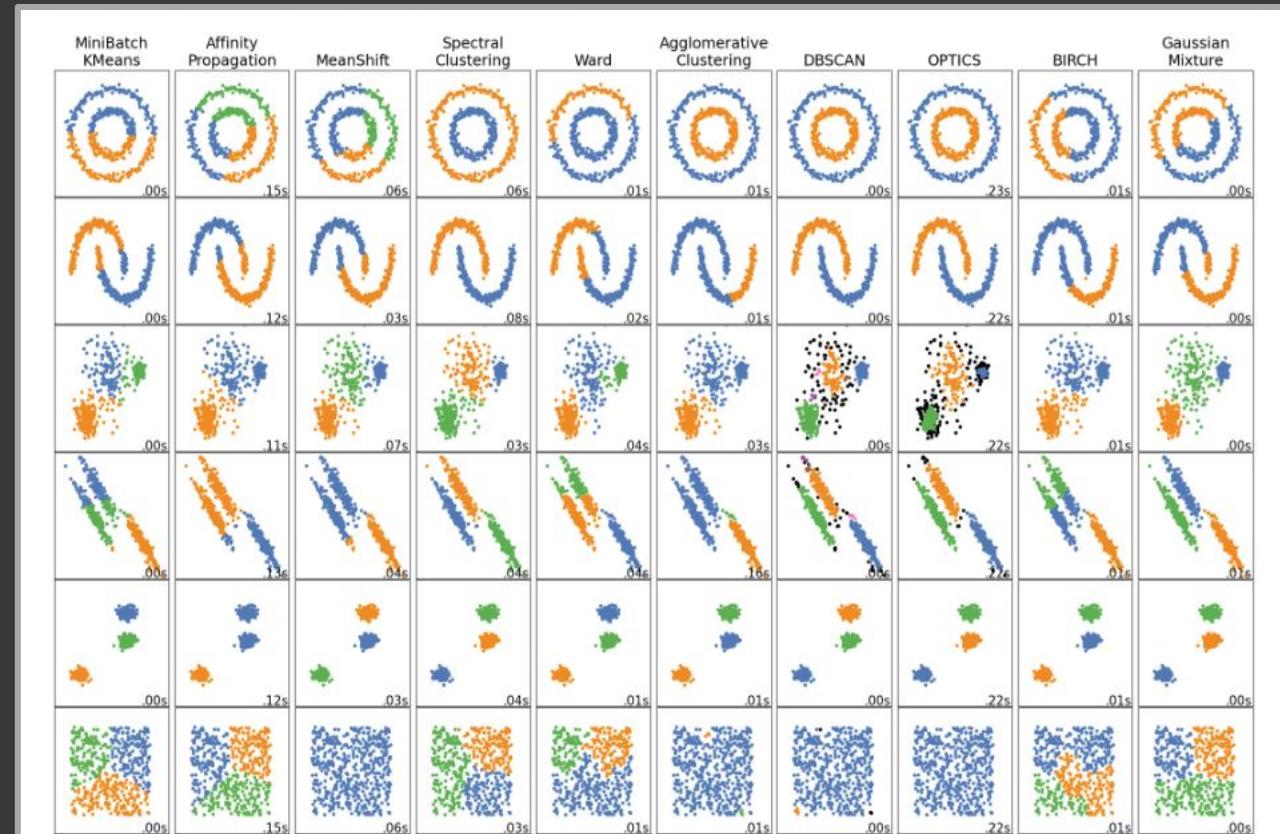
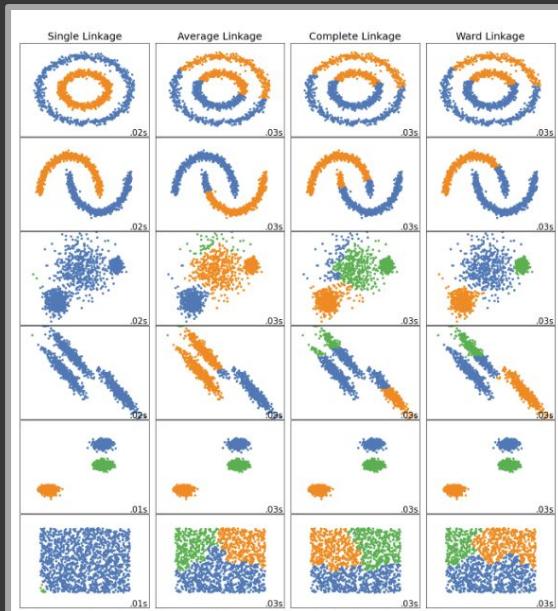


The solution: calculate the local ground average z-value (height above ground level) and reclassify the tile

# OUR APPROACH

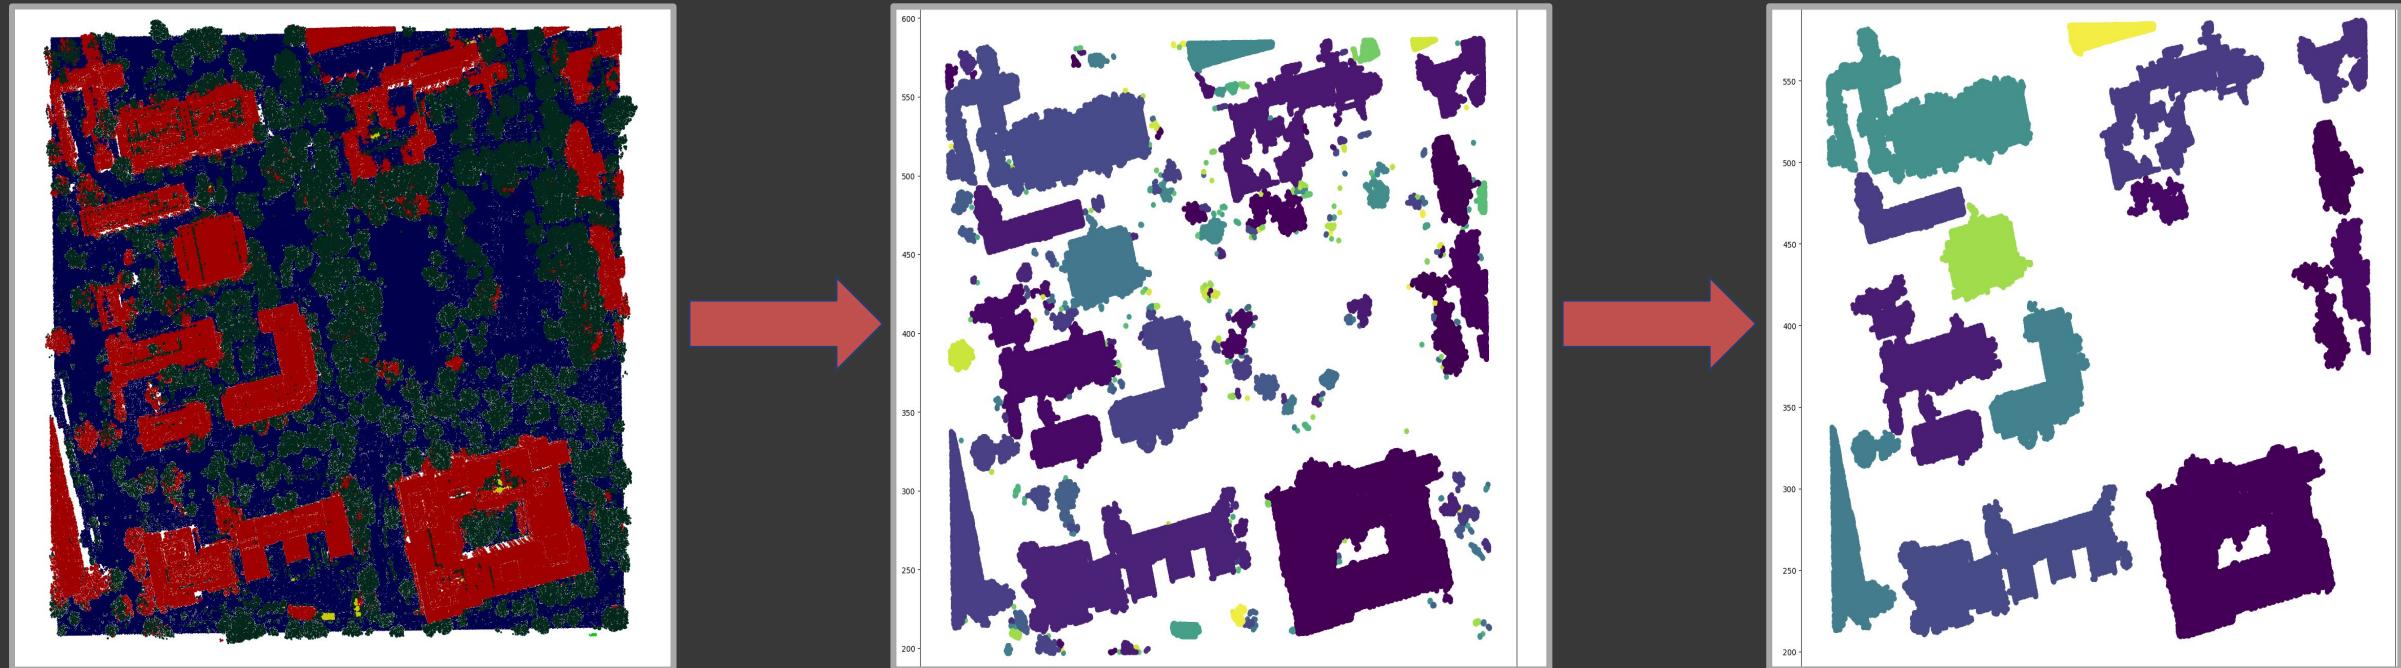
## Segmentation & Clustering

- Run clustering algorithm on points each class separately
- Multiple segmentation methods were explored
- Most of the clustering algorithms use the bottom up approach which are very memory intensive (requires at least 60GBs - 100GBs of RAM)
- Agglomerative Clustering with Single linkage using the distance threshold



# OUR APPROACH

## Clustering and Noise

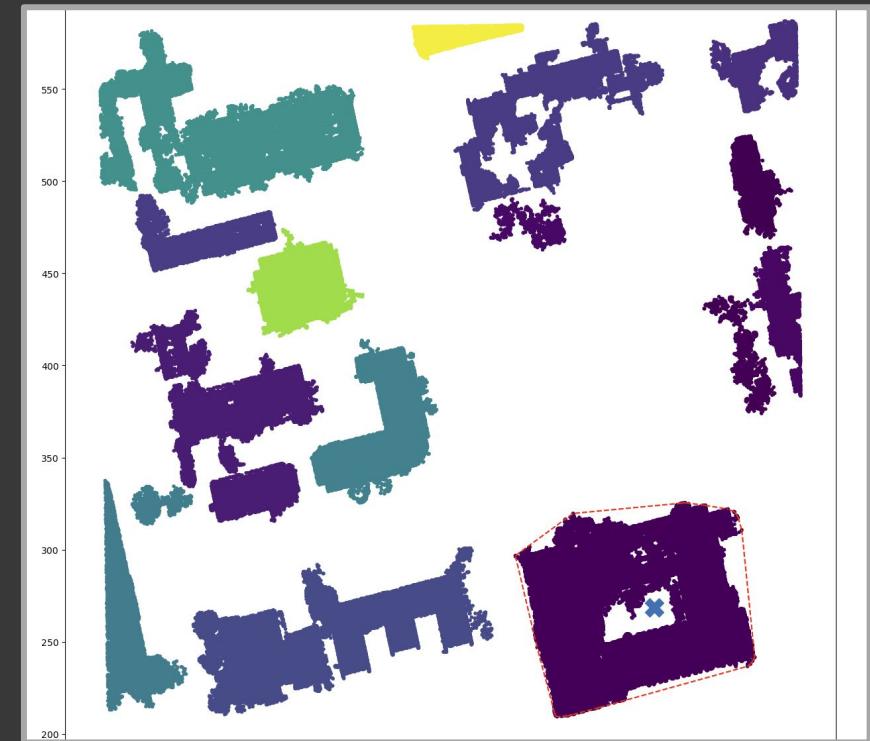
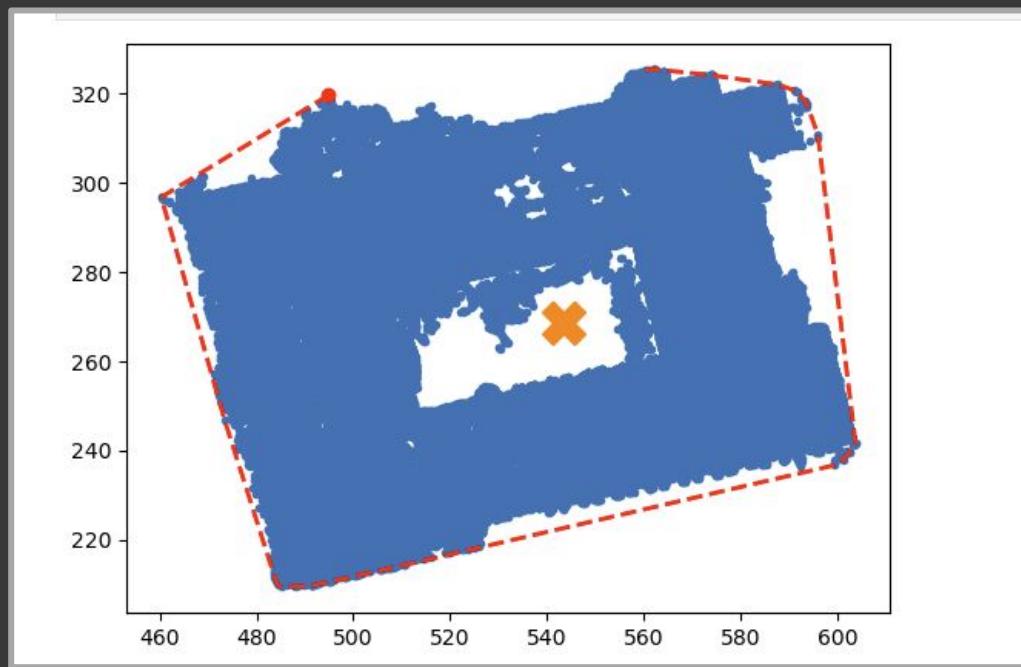


- We observed an overclassification of buildings. We knew we had to remove some of the noise
- To combat this overclassification, we used a custom function to remove clusters of point clouds made up of less than 1000 points
- In this example, initial clustering returned 312 clusters. After removing the noise, we ended up with 16 clusters

# OUR APPROACH

## Convex Hull

- A convex hull bounds the clusters with a bounding box, using the set of it's extreme points
- We calculated the 'centroid' of the bounding box using the mean of all x-values and y-values
- This point was then used as a pinpoint for the object, which provides a faster database query, compared to using the vertices of the bounding box



# OUR APPROACH

Our first model is successful for classifying static objects. However, we needed a solution to identify dynamic objects

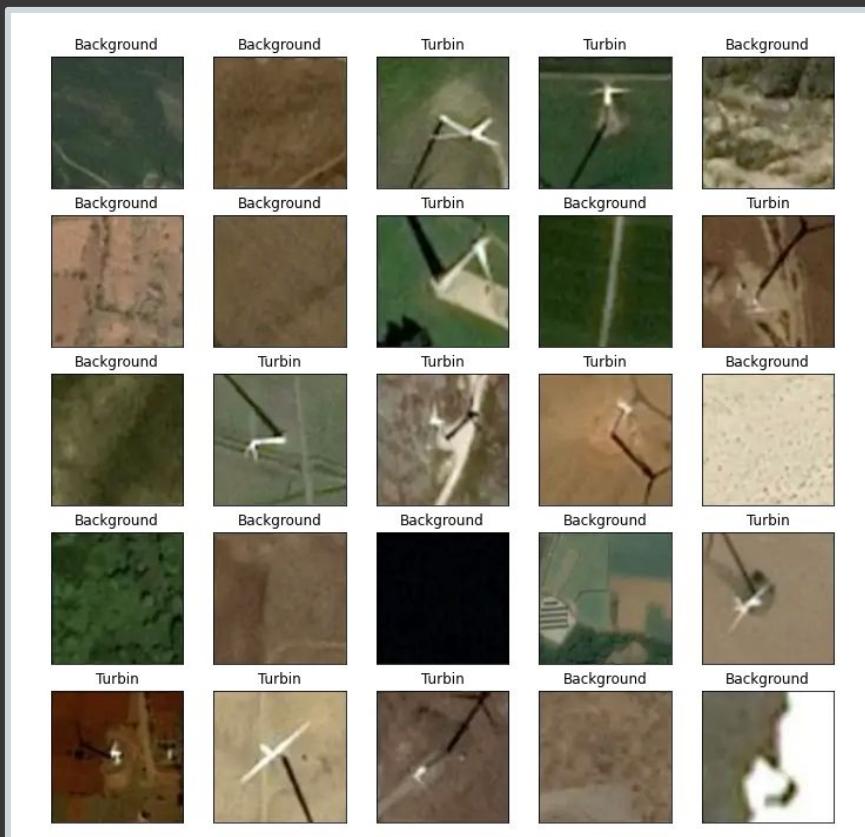


# OUR APPROACH

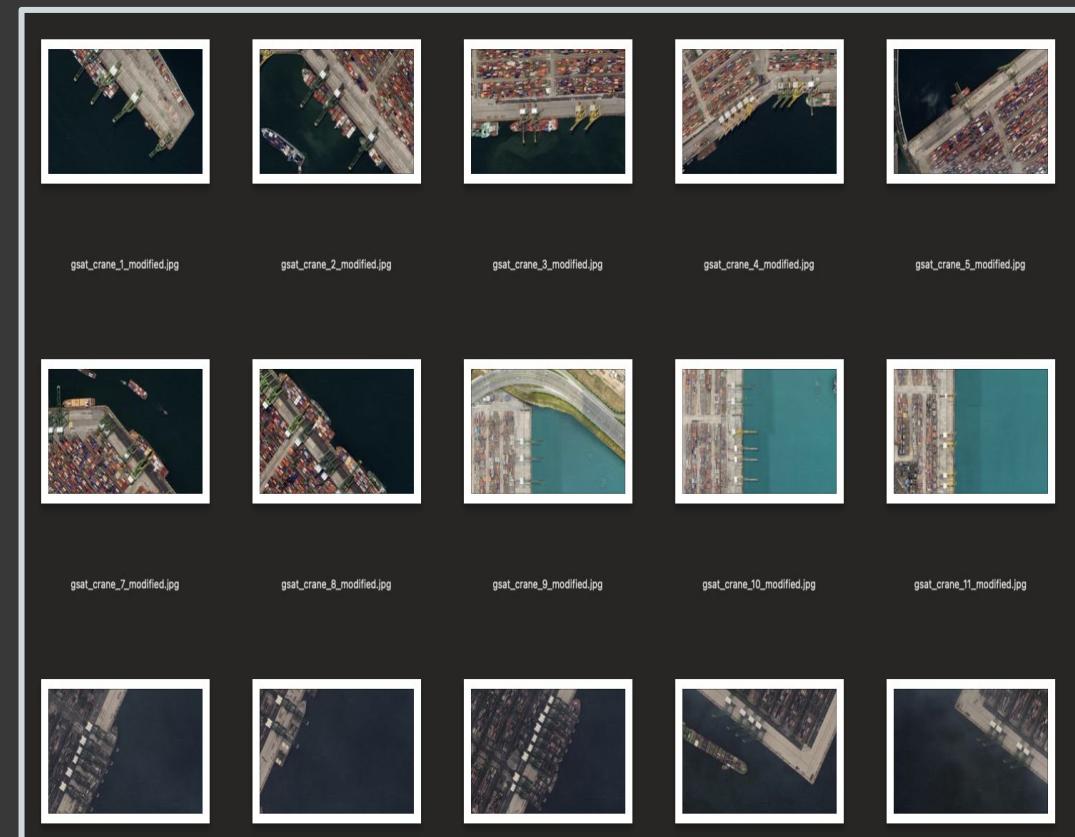
## Classification

Satellite Image Classification to handle dynamic obstacles

Dataset: Airbus SPOT wind turbines patches



Dataset: IEEE DataPort 0.5-m resolution of cranes



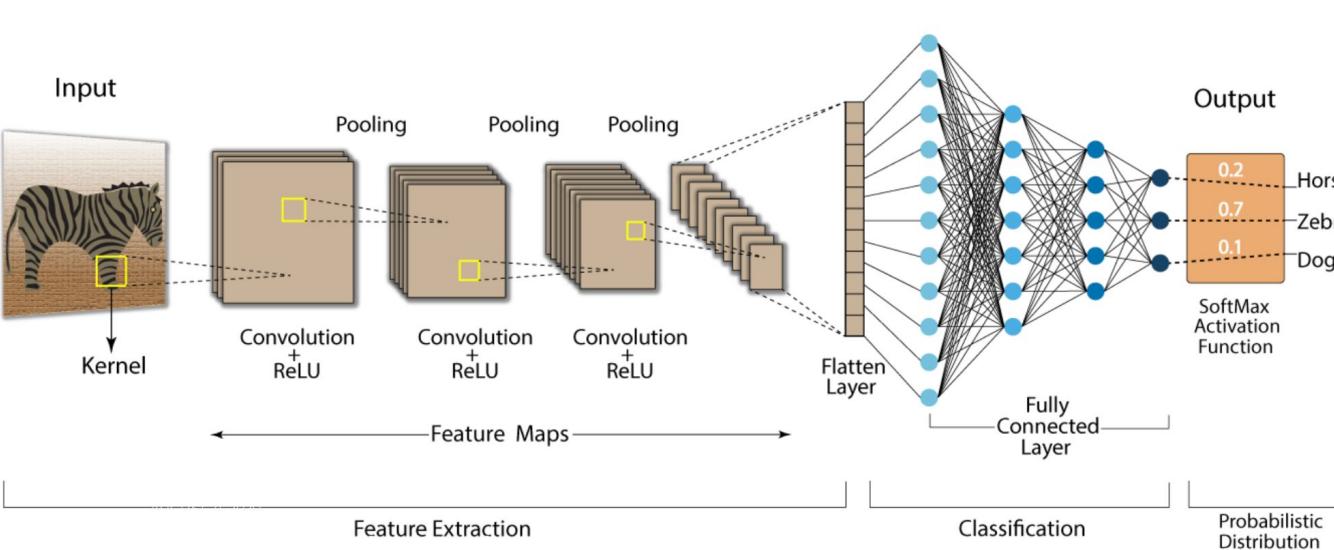
# OUR APPROACH



## Second Model - Convolution Neural Network (CNN)

- Three Convolutional Layers (32, 32, 64 filters)
- 128 hidden neurons on dense/fully connected layer
- Total of 7.5 million parameters
- Transformation & Image Augmentation using Keras: (zoom, flip, rescale, shear etc.)
- Achieved 99.7% accuracy

**Convolution Neural Network (CNN)**



cnn.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 32)	0
conv2d_2 (Conv2D)	(None, 60, 60, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 64)	0
flatten (Flatten)	(None, 57600)	0
dense (Dense)	(None, 128)	7372928
dense_1 (Dense)	(None, 1)	129

Total params: 7,401,697  
Trainable params: 7,401,697  
Non-trainable params: 0

# OUR APPROACH

## Scalability Challenges

### Challenges:

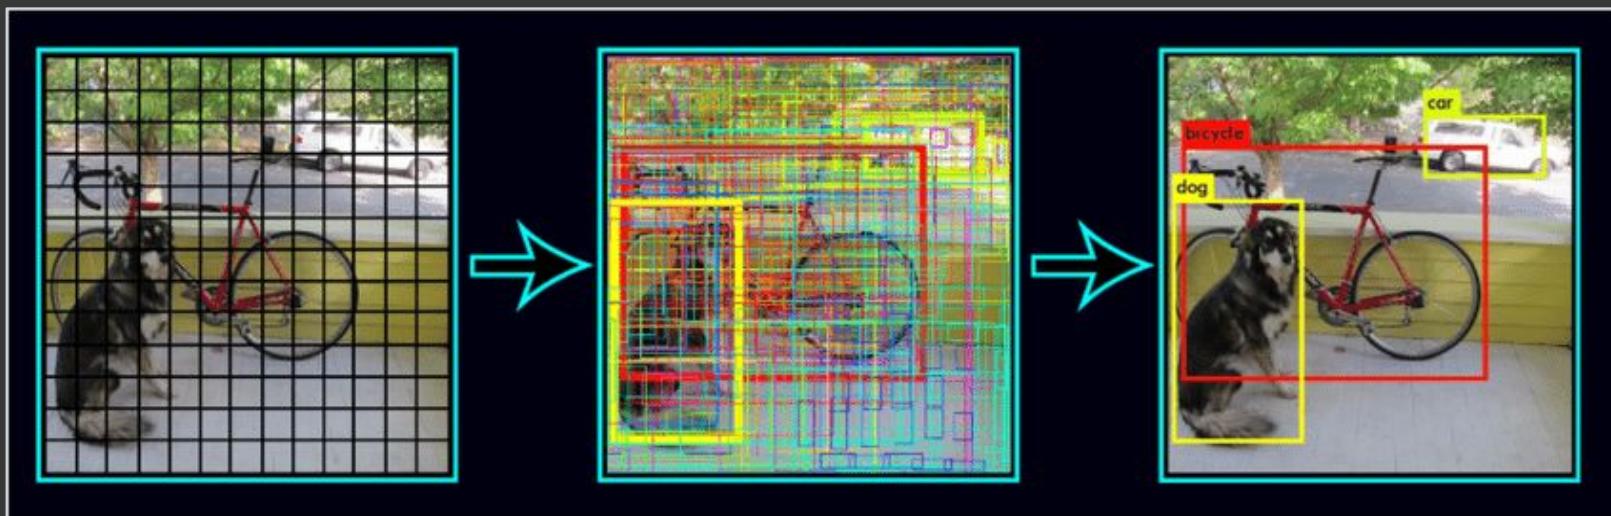
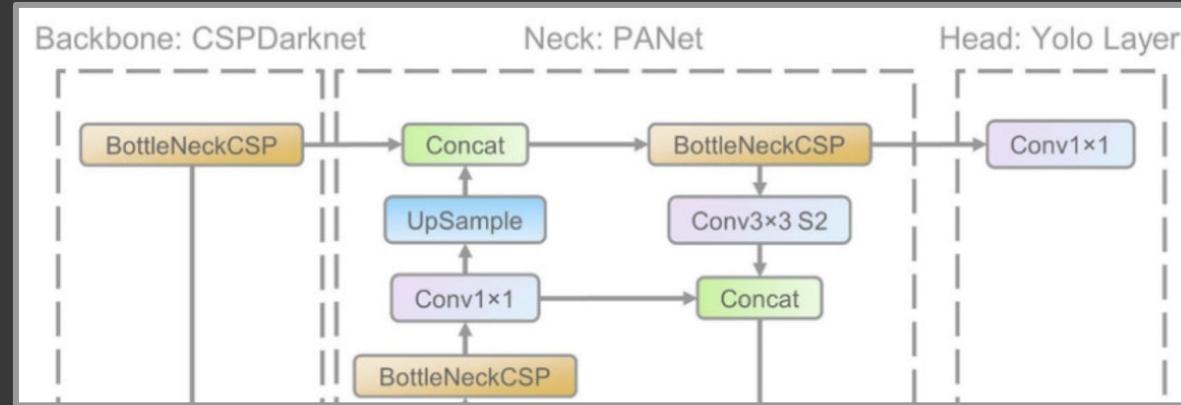
- Unable to handle multiple objects
- Unable to provide location
- Object must be centered within image
- Object must be zoomed in
- Computational and monetarily taxing to scale up



# OUR APPROACH

## Object Detection with Bounding Box: YOLO

- YOLO algorithm (You only look once) single stage detector
- Object detection algorithm that divides images into a grid system
- Each cell in the grid is responsible for detecting objects within itself



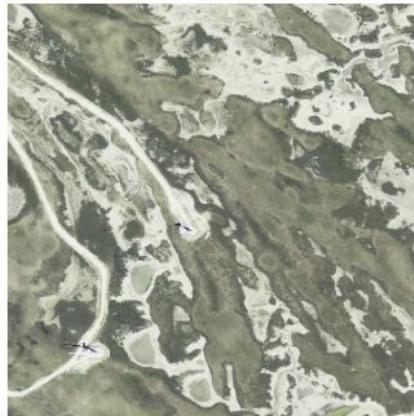
# OUR APPROACH

## Assemble the Dataset

### Diversity Between States



Maine



Texas



Idaho



California

- NAIP Power Plant Aerial Imagery Dataset
- 4454 Overhead images across the US from California, the Midwest and Texas
- Has variations in land types, wind turbine shapes and sizes across states
- By including variation in features like geography, wind turbine size, and shadow size in the training dataset, the model can generalize more broadly

Figure: Images are taken from National Agricultural Imagery Program (NAIP) Power Plant Aerial Imagery Dataset.

# OUR APPROACH

Roboflow

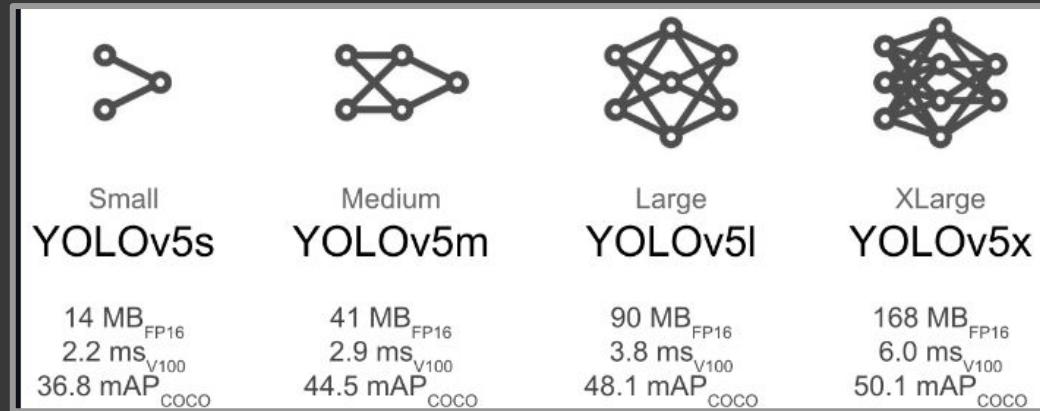
Using Roboflow:

- Convert an existing dataset to YOLOv5 format
- Upload raw images and annotate them in Roboflow with Roboflow Annotate
- About 800 images were annotated for model training and validation
- Augmentations (Resize, zoom, flip and rotation) were applied



# OUR APPROACH

## Model Training



- We used a Pytorch implementation of YOLOv5 Model to make use of the Roboflow dataset
- We chose the generic YOLOV5x (XLarge) checkpoint pretrained on Common Object in Context Dataset (COCO)
- 150 Epochs
- 322 layers
- >8 million parameters

150 epochs completed in 2.374 hours.

Optimizer stripped from runs/train/exp3/weights/last.pt, 173.1MB

Optimizer stripped from runs/train/exp3/weights/best.pt, 173.1MB

Validating runs/train/exp3/weights/best.pt...

Fusing layers...

Model summary: 322 layers, 86173414 parameters, 0 gradients, 203.8 GFLOPS

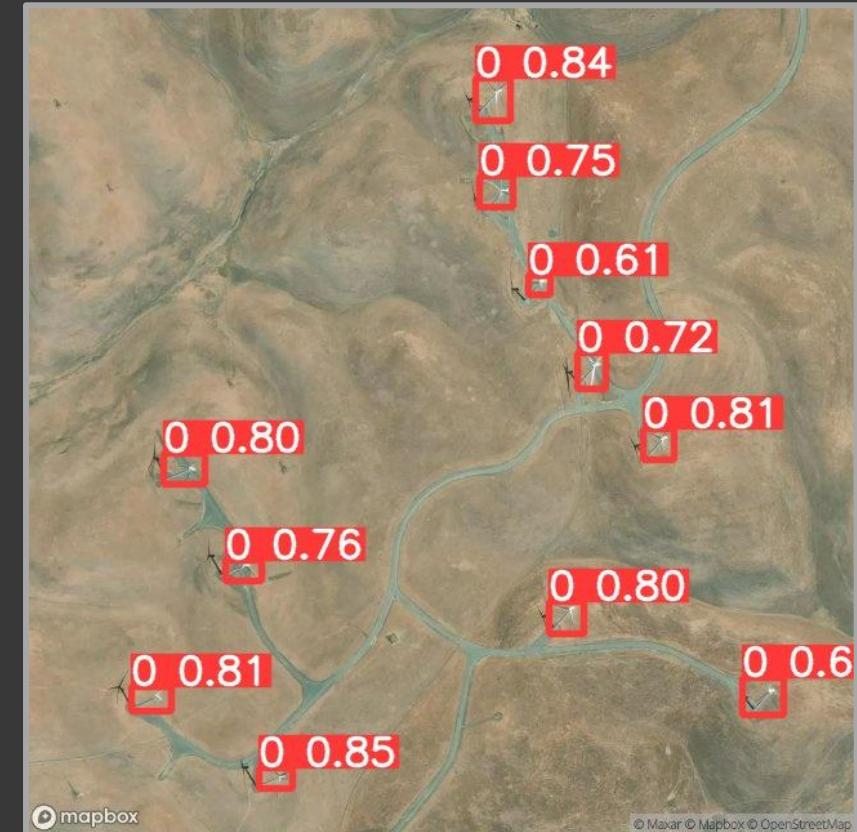
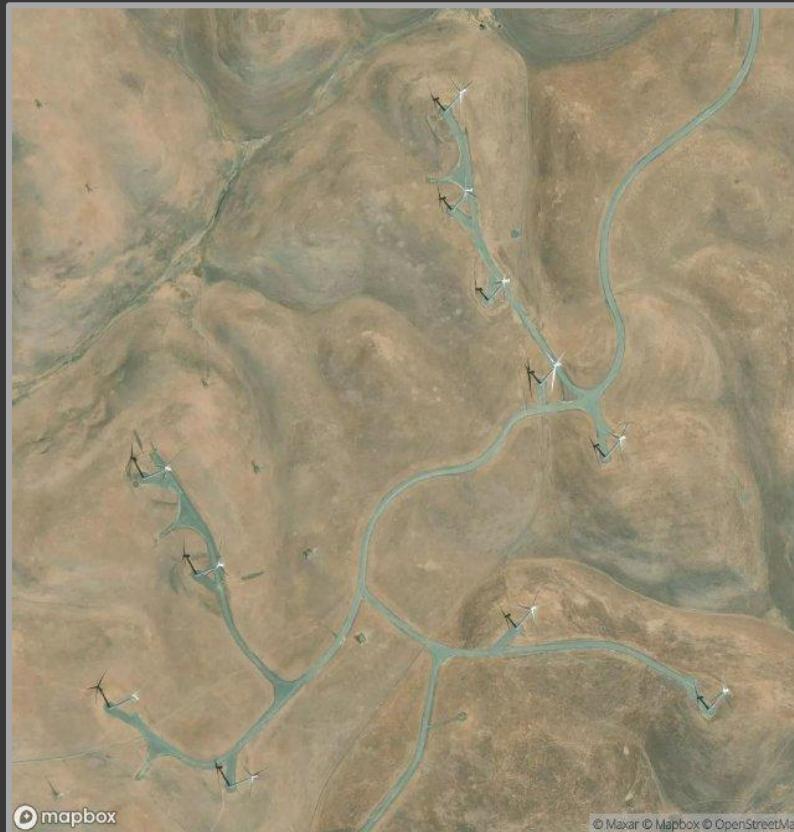
Class	Images	Instances	P	R	mAP50	mAP50-95:	100% 3/3 [00:02<00:00, 1.04it/s]
all	70	373	0.731	0.708	0.685	0.25	

Results saved to runs/train/exp3

# OUR APPROACH

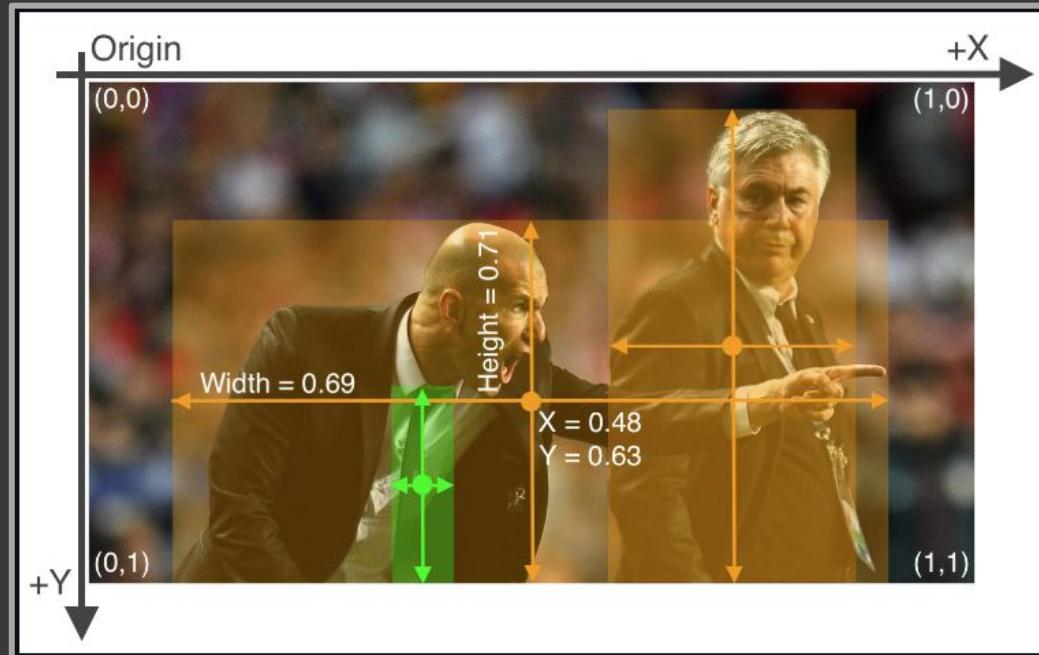
## Model Results

- Tested in a real world dynamic setting
- Near perfect accuracy
- Baseline parameters were reduced due to safety concerns



# OUR APPROACH

## Bounding Box to Geo Location



   zidane.txt ▾

0	0.481719	0.634028	0.690625	0.713278
0	0.741094	0.524306	0.314750	0.933389
27	0.364844	0.795833	0.078125	0.400000

Parameters

access token from your account page  
`pk.eyJ1Ijoic3RldmVuenp6Iiw1YSI6ImNs`

style  
Mapbox Satellite

Use my own public style

width in pixels height in pixels  
640 640

center  bounding box  auto

min longitude max longitude  
-121.7997 -121.7845

min latitude max latitude  
38.0742 38.0862

Optional parameters

padding add extra space inside the map  
50,40,30,10

overlays add a marker, path, or GeoJSON  
+ add overlay

@2x render map at 2x scale

style parameters  
change map style at render time  
add a layer set a filter

Request URL

[https://api.mapbox.com/styles/v1/mapbox/satellite-v9/static/-121.7997,38.0742,-121.7845,38.0862/640x64?access\\_token=pk.eyJ1Ijoic3RldmVuenp6Iiw1YSI6ImNs](https://api.mapbox.com/styles/v1/mapbox/satellite-v9/static/-121.7997,38.0742,-121.7845,38.0862/640x64?access_token=pk.eyJ1Ijoic3RldmVuenp6Iiw1YSI6ImNs)



A satellite map showing a coastal area with a large body of water and land. A rectangular bounding box is overlaid on the map, indicating a specific geographic area. The Mapbox logo is visible in the bottom left corner of the map area.

Each row → Class, class x\_center, y\_center, width, height

# OUR RESULTS

## SQL Database

Object	HAGL	Center	Vertices
0 Building(s)	10.79	[38.542552457073164,-121.75018286955013]	[38.542589003007436, -121.75059566322228],[38.542589003007436, -121.75059566322228]
1 Building(s)	28.67	[38.54224087158561,-121.74948327765163]	[38.54198864246644, -121.75031621258651],[38.54198864246644, -121.75031621258651]
2 Building(s)	16.89	[38.542405173389525,-121.74812607876285]	[38.54250342930847, -121.7485321144085],[38.54250342930847, -121.7485321144085]
3 Building(s)	31.25	[38.542138962823564,-121.75192774428984]	[38.542291635885064, -121.75102776608105],[38.542291635885064, -121.75102776608105]
4 Building(s)	13.10	[38.541877287834126,-121.74831633755196]	[38.541767534024586, -121.74800868018991],[38.541767534024586, -121.74800868018991]
5 Building(s)	18.87	[38.541149374871544,-121.74806758924065]	[38.54138842665915, -121.74826240589809],[38.54138842665915, -121.74826240589809]
6 Building(s)	20.21	[38.54098145020309,-121.74838440926165]	[38.54070194020875, -121.74820488813141],[38.54070194020875, -121.74820488813141]
7 Building(s)	21.11	[38.539740390475984,-121.74898894191779]	[38.54020549320194, -121.74953724169754],[38.54020549320194, -121.74953724169754]
8 Building(s)	29.48	[38.54162312730218,-121.75213536398749]	[38.541654393264814, -121.75251207130829],[38.541654393264814, -121.75251207130829]
9 Building(s)	16.42	[38.54162414107391,-121.74978670225218]	[38.54170032339375, -121.75004249482232],[38.54170032339375, -121.75004249482232]
10 Building(s)	12.96	[38.541345884443864,-121.75138283328235]	[38.54161502949612, -121.75150772970817],[38.54161502949612, -121.75150772970817]
11 Building(s)	16.70	[38.54073896608736,-121.75189648280059]	[38.54019055891056, -121.75159133412446],[38.54019055891056, -121.75159133412446]
12 Building(s)	15.50	[38.540588390193356,-121.75086420944267]	[38.54049801890671, -121.75132796884193],[38.54049801890671, -121.75132796884193]
13 Building(s)	17.10	[38.53971263266667,-121.75262723815835]	[38.539410481866135, -121.75223897733719],[38.539410481866135, -121.75223897733719]
14 Building(s)	15.55	[38.54027602574538,-121.75238704727289]	[38.540281701972944, -121.75256687983682],[38.540281701972944, -121.75256687983682]
15 Building(s)	20.60	[38.53959957908259,-121.75125555976568]	[38.53964815933765, -121.75010566891191],[38.53964815933765, -121.75010566891191]

### Object:

- Classified obstacle

### HAGL (Height Above Ground Level):

- Calculated (static)
- Assigned fixed value (dynamic)

### Center:

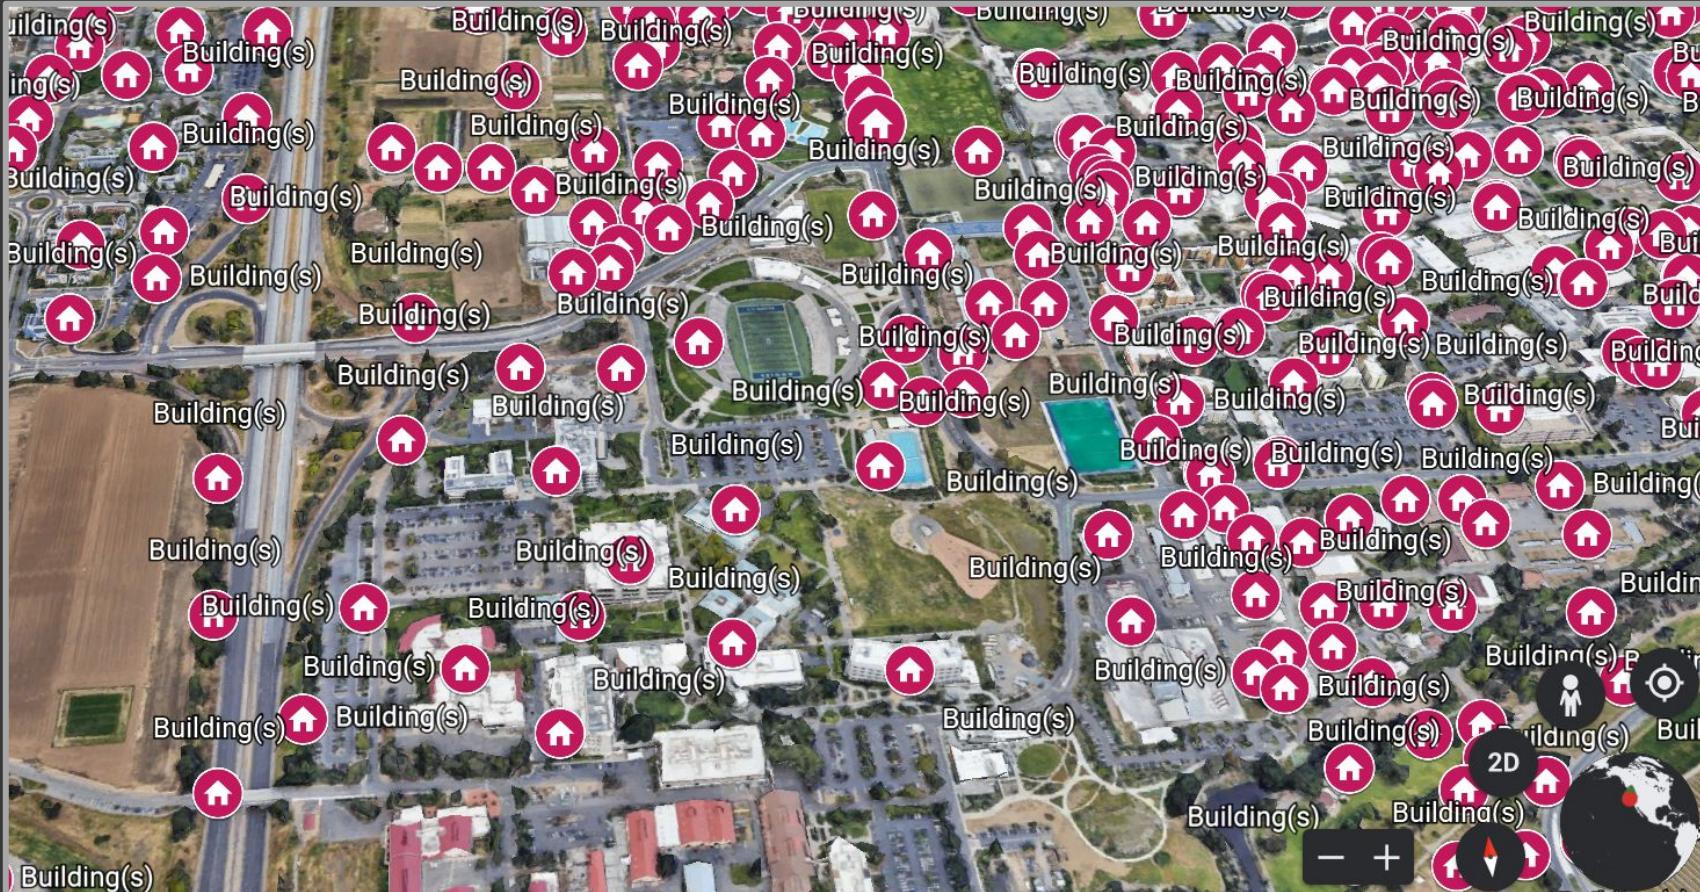
- Centroid of convex hull (static)
- Width and height relative to image (dynamic)

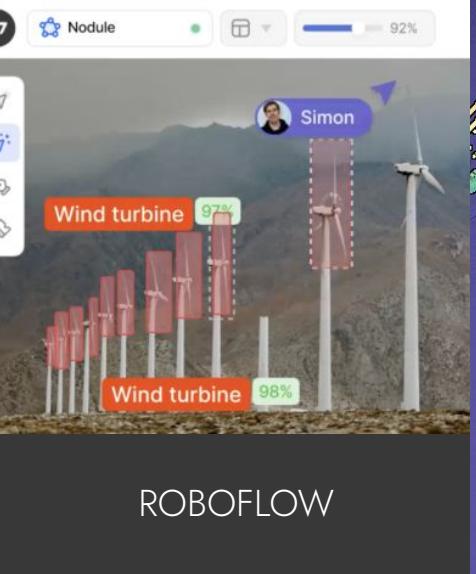
### Vertices:

- Outer extreme points of convex hull

# Live Demo

[Link](#)





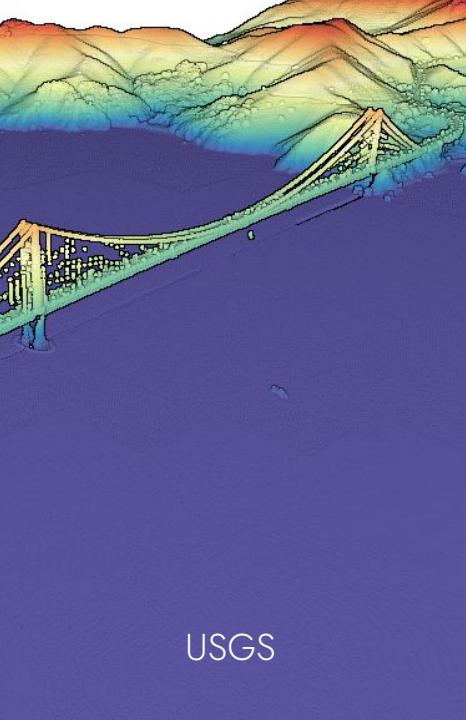
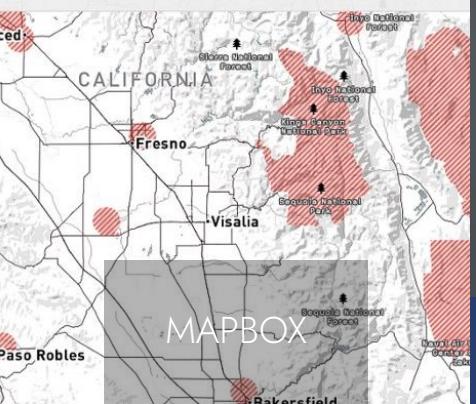
DON'T FLY DRONES HERE

Small unmanned aircrafts are constantly in flux. This map represents a live location of areas where it is not recommended to fly. If you'd like to add a site or remove one, please open an issue in the [GitHub repository](#).

NO FLY ZONES

Major Airports \* U.S. Military Bases U.S. National Parks

Radius around medium to large sized airports



USGS



## Tools

# Tech Stack

## Tools

Roboflow (Annotation/labeling Images)  
TensorBoard (performance metrics)  
MapBox API (Satellite images)  
USGS API (Lidar data)  
GCPVM (Virtual machine),  
Ubuntu Desktop (Linux GUI),  
Chrome Remote Desktop (VM remote access),  
Tmux (Terminal Multiplexer),  
Github (Collaboration & version Control),  
JIRA (agile development & Project management),  
Confluence (Documentations),



# Tech Stack

## Libraries

### Framework

Tensorflow

Pytorch

### Libraries

Keras (Deeplearning API)

trimesh (Dealing with meshes from binary/ASCII)

plyfile (Reading and writing ASCII and binary PLY file)

laspy (Lidar pointCloud io)

utm (Bidirectional UTM-WGS84 converter)

open3d (Data processing & visualization)

sqlite3 (Interaction with SQL database)

Common Lib

Pandas, matplotlib, sklearn, sys, os, cv2, numpy, PIL, traceback, scipy, glob

### Database & DBMS:

SQL & mySQL



# FUTURE WORK

What we still want to explore

CI/CD pipeline (constant updates of the Database)

Faster and more specialized database in geospatial data: PostGIS, Snowflake, Redis

Structure from motion (another object classification method)

Optimization (using classified objects to determine risk and path)

Improve our models by training on more annotated data (more 'good' data is more important than better algorithms)

Explore a more efficient framework for point cloud classification (downsampling & TONIC)

# Thank You

Q & A

